

व्यायाम

Vyāyāma

On-device AI Form Coach for Snapdragon

PROJECT BIBLE

Problem statement FitSense — Real-Time Exercise Detection & Recognition

Target hardware Qualcomm QIDK · Snapdragon 8 Elite (Hexagon V79) · Android

Version 1.1 · **Self-audit** 94.2 / 100 over 3 iterations

0. Front matter

v1.1 — DEVICE CONFIRMED (organizer erratum on the PS): the venue device is a **QIDK with Snapdragon 8 Elite (SM8750, Hexagon V79), running Android** — *not* the QCS6490/RB3 board the PS image showed. This **resolves landmines L1/L2/L4/L5** (\$22): one known flagship target (top-tier NPU), the reference **VisionSolution4 is validated on 8 Elite** (its `.d1c` + `skel .so` are the right ones), the `<80 ms` budget is trivial, and the developer-kit Android build should permit `adb root / setenforce 0`. ADB ships with Android Studio platform-tools. The board / V68 / USB-cam-HDMI notes below are now dead-path insurance only.

0.1 How to use this document — the three governing rules

1. **Default to the bible.** When in doubt, what this document says wins. It is grounded in the official QIDK Quick Start PDF and verified research — not memory.
2. **Each engineer owns their section.** Three tracks (A "Metal", B "Brain", C "Glue→App"). Even if one person does several, keep the **module boundaries clean** — they exist so a failure in one track never cascades into another.
3. **Independence is non-negotiable.** No module tightly couples to another. **Every interface has a mock**, and the app must build and run end-to-end with **any subset of real implementations**. This is what lets the team work in parallel and what guarantees a demo even on partial completion.

0.2 The non-negotiable engineering vows (read before writing a line of code)

- **No invented numbers.** Every latency / FPS / accuracy figure is either measured on-device or explicitly labelled "not yet measured." A fabricated number is a fireable offence in this project, because Qualcomm engineers judge this track and will probe it.
- **Honest `backend()` always.** The app reports the compute backend it *actually reached* (NPU / GPU / CPU). It never claims NPU when it silently fell back. The honesty is the credibility.
- **Zero allocation in the per-frame path.** Pre-allocate every buffer at startup and reuse it. No `FloatArray( ... )`, `ByteArray( ... )`, `Bitmap.create...`, or per-frame JNI object marshalling in the analyzer loop.
- **The bible wins.** See rule 1.

0.3 Changelog

- **v1.0** — first full draft from the approved plan. (Iterations 1–3 hardening recorded in Appendix D.)

1. Executive summary & win thesis

Vyāyāma is a fully on-device fitness **form coach**. You stand in front of a Snapdragon device, start exercising, and the app — running entirely on the chip, no cloud, no internet — detects *that* you're exercising, recognises *which* of five exercises you're doing, **counts your reps, scores each rep's form 0–100**, and speaks **one corrective cue at a time** ("go deeper", "knees out", "hips up"). It also shows a live **CPU-vs-NPU latency comparison** proving the work runs on the Hexagon NPU.

The strategic reality that defines our entire approach: Qualcomm handed every team in this track a working reference app — [VisionSolution4-PoseEstimation](#) — that already runs **person-detection + pose estimation on the Hexagon NPU** with a built-in CPU/GPU/DSP runtime toggle. So pose-on-NPU is **table stakes**, not a differentiator. The room will be full of teams that reskin the sample, and a Qualcomm engineer will see through that in thirty seconds.

Win thesis: *The reference ends at keypoints; our product begins there.* We fork the reference, keep its proven native inference core verbatim, and build the layer Qualcomm's sample stops at — a transparent, explainable **biomechanics + coaching engine** on top of the 17 keypoints, wrapped in an honest, live CPU-vs-NPU demonstration. We don't compete on the part everyone was given; we go deep on the part nobody was given, and we are radically honest about the seam between the two.

Four reinforcing positioning angles:

1. **Depth over the reference, not breadth around it** — the biomechanics layer is the hard, un-given half.
2. **Honesty as a moat** — measured-only numbers + honest backend labelling pre-empt the judge's probe.
3. **A demo that is a controlled experiment** — the CPU↔NPU toggle is a *result that changes live*, not a claim.
4. **Genuinely useful + private** — offline form coaching with no subscription and no video leaving the device.

Hero use-case: AI Form Coach with a hint of gamification (streaks / personal-best / per-session grade). **First exercise set (Core 5, bodyweight, single-person, front-camera):** squat, push-up, lunge, bicep curl, jumping jack.

2. The product in plain language (for a non-technical judge)

You open Vyāyāma and prop the phone up. You start doing squats. Instantly a stick-figure skeleton tracks your body, a big label says **SQUAT**, and a counter ticks **1 ... 2 ... 3** as you complete each rep. Next to each rep a score flashes — **92, 88, 95** — how good that rep's form was. When you do a sloppy squat (not deep enough, knees caving in), the score drops and the app shows one clear instruction: **"Go deeper."** Not a wall of warnings — one fix at a time, like a real coach. At the end you get a summary: 20 reps, average form 90, your most common mistake, and a streak badge.

There's also a small panel showing how fast the AI is running — and a button that switches it between the phone's **NPU** (the dedicated AI chip) and its **CPU**. Flip it and you see the speed drop on CPU and jump back up on NPU. That's the proof the heavy lifting really runs on Qualcomm's AI hardware.

Everything happens **on the device**. No video is uploaded. No account. No internet needed.

3. The "not a reskin" chapter (read this to a Qualcomm judge first)

This is the most important section in the document, because "you just reskinned the sample" is the one objection that can sink an otherwise-strong project in this track. We disarm it before it's raised.

The PDF's Chapter 6 explicitly invites teams to go beyond the default model ("Participants are encouraged to explore better open-source models... More keypoints = better downstream results for gesture/pose recognition"). Building beyond the reference is *responding to the brief*, not going off-script. We make that explicit and honest.

The given-vs-built ledger (goes on a slide and in the deck)

What Qualcomm gave (the reference VisionSolution4)

What Vyāyāma built on top

YoloNAS person-detection on the NPU	Exercise-state gating — detecting <i>when</i> someone is mid-exercise vs idle/resting
HRNet 17-keypoint pose on the NPU (SNPE C++)	A joint-angle feature layer — per-frame biomechanical features (knee/hip/elbow/shoulder angles, symmetry, depth ratios, torso lean) with smoothing + confidence gating
Runtime toggle CPU/GPU/DSP + fallback ladder	An instrumented, demoable CPU-vs-NPU comparison — live per-frame latency in the UI, warm-up-corrected, honest backend-reached labelling
<code>.dlc</code> for Hexagon (V79-targeted)	Per-device <code>.dlc</code> discipline (V68-class for QCS6490) + a device-agnostic runtime story
Skeleton rendering	The Form Coach — per-exercise rep counting (state machine), per-rep form scoring, one-cue-at-a-time corrective coaching, session summary, light gamification
(nothing)	Exercise recognition for the Core 5 — rule-based, with an optional learned temporal head as the scale path

The one-line weapon for Q&A: *"The reference answers 'where are the joints?' We answer 'is your form going to hurt you, and how do you fix it?' — on the device, with no cloud."*

Why this wins respect: it shows we *understood* what we were handed (most reskin teams can't articulate it), it reframes our work as the high-value half, and it makes our honesty an asset. We **under**-claim the given part and **over**-deliver the built part — the opposite of the teams that will get caught.

Concretely, vs a typical reskin team: a reskin ships the reference's skeleton with a logo and maybe a hard-coded "squat counter" on one joint threshold (double-counts on jitter, no form feedback, no "is-exercising" gate, claims "NPU" with no measurement). Vyāyāma ships a confidence-gated feature layer, a hysteresis rep FSM with partial-rep handling, *corrective* per-fault cues with a 0–100 score, an honest warm-up-corrected CPU-vs-NPU panel, and a device-agnostic runtime that labels its true backend. The gap is visible in ten seconds of demo and survives every follow-up question.

4. Problem-statement → solution mapping (proof every requirement is solved)

Every line of the 3-part FitSense brief mapped to the component that satisfies it and the track that owns it.

#	PS requirement	Solved by	Bible \$	Track
P1.1	Detect when a person is exercising	"is-exercising" gate in <code>ExerciseClassifier</code> (idle vs active from pose motion energy)	\$9	B
P1.2	Identify which exercise	<code>RuleExerciseClassifier</code> (Core 5) + optional learned head; honest UNKNOWN	\$9	B
P1.3	Run fully on-device , no cloud/internet	No network permission in the prod path; all compute on the SoC	\$6, \$20	A/C
P2.1	Convert model to <code>.dlc</code> for the Hexagon NPU	DLC generation (<code>GenerateDLC.ipynb</code>) or AI Hub Path B per device; runs on HTP via SNPE	\$10	A
P2.2	Target <80 ms/frame	Labelled per-frame budget; HTP BURST + zero-copy; detect-every-N / single-stage levers; measured-only	\$11	A
P2.3	Live CPU-vs-NPU comparison	<code>Benchmark</code> + reference runtime toggle; warm-up-corrected p50/p90/FPS; honest backend-reached; on-stage NPU→CPU jump	\$10	A/C
P3.1	Build a useful app	AI Form Coach: rep counting + per-rep form score + one-cue coaching + session summary + light gamification	\$12–15	C/B
P3.2	Help someone work out better	Actionable corrective cues (depth, valgus, hip-sag, swing, symmetry); private, offline, no subscription	\$13	B

If a judge asks "does it do X?", point at the row. Nothing is hand-waved.


5. Judging criteria & our per-criterion play

Likely criteria for this track: innovation, technical implementation, impact, feasibility, scalability, presentation. What earns a top score, and where we are honestly weak today (pre-build).

- **Innovation — strong.** The biomechanics layer is non-trivial: per-exercise FSM rep counting, *corrective* (not binary) form cues, a per-rep 0–100 score, plus the optional learned head. *Weak spot*: "rep counting from angles" is a known idea — so we make the **corrective cue engine** the star (most apps only count) and avoid ever claiming the *pose model* is our innovation (it's Qualcomm's).
- **Technical implementation — strong, with one real risk.** Clean layered architecture, swappable pose source, unit-tested math, and a working live CPU-vs-NPU switch. *Weak spot*: getting the NPU path running on the *actual* venue device (right Hexagon arch, runtime libs, root/SELinux). De-risked by keeping the reference core verbatim, proving NPU on ≥2 devices in prep, and making the CPU path genuinely good.
- **Impact — strong if framed right.** Private, offline, personal-trainer-grade form feedback; injury prevention; accessibility (a coach is expensive, this is free + offline). *Weak spot*: "fitness app" sounds unserious to a hardware judge — so lead with **privacy + accessibility + injury-prevention**.
- **Feasibility — good, must stay disciplined.** Demonstrably *done* beats promised: the mock-mode app runs with zero hardware, so even partial success demos. *Weak spot*: solo/small-team + venue-gated NPU + per-device DLC — front-loaded into prep (see §23).
- **Scalability — strong, with an honest ceiling.** Adding an exercise = adding one analyzer module (no retraining for the rule path); pose source is model-agnostic; the runtime ladder is device-agnostic. *Weak spot*: rule-based form is per-exercise hand-tuning — named openly, with the learned head as the scale story.

- **Presentation — high ceiling, solo risk.** One unmistakable wow beat (live latency drop / a bad rep getting corrected), a clean projector-legible overlay, and the honesty beat in the first 60 seconds. *Weak spot:* one person narrating + device-juggling — mitigated by rehearsal + a backup video.

6. System architecture



Two data paths are drawn explicitly: the **coaching path** (camera → pose → intelligence → CoachState → UI) and the **benchmark path** (PoseEngine reports per-inference latency + backend → Benchmark → the live A/B panel). The UI is a pure function of the latest `CoachState`; it never blocks inference.

7. Module boundaries & independence (the interface-with-mock discipline)

A small set of Kotlin interfaces, each with a **real** impl in its module and a **mock**, wired by a tiny hand-rolled DI container (no Hilt/Koin — keep hackathon build times down). The app boots and demos with **any subset mocked**.

Interface	Real impl	Mock	What the mock unblocks
PoseEngine	SnpePoseEngine (JNI→SNPE)	MockPoseEngine (replays recorded keypoint CSV @30fps)	Keystone — entire app runs with no device/NPU/camera
CameraEngine	RealCameraEngine (CameraX YUV, zero-alloc)	MockCameraEngine (no-op; mock pose drives frames)	UI/intelligence dev on emulator
FeatureExtractor	RealFeatureExtractor (pure Kotlin)	— (pure fn, unit-tested directly)	n/a
ExerciseClassifier	RuleExerciseClassifier / LearnedExerciseClassifier / FusedExerciseClassifier	MockExerciseClassifier (fixed type)	rep/form/UI dev against a known exercise
RepCounter	StateMachineRepCounter	MockRepCounter (timer-increment)	UI dev
FormAnalyzer	RuleFormAnalyzer	MockFormAnalyzer (canned cues/score)	UI polish
Benchmark	RealBenchmark (EMA/p50/p90/FPS)	MockBenchmark (labelled fake numbers)	latency panel renders in mock mode

The decisive subsets:

Scenario	Real	Mocked	Capability
Pure UI dev (laptop emulator)	—	all	build/polish entire UI, overlay, gamification, latency panel
Intelligence dev (any phone)	FeatureExtractor, Classifier, RepCounter, FormAnalyzer	PoseEngine (CSV replay), Camera, Benchmark	develop + unit-test the whole brain on recorded sessions; no SNPE/ .dlc
Native bring-up (venue device)	PoseEngine (real SNPE), Camera, Benchmark	intelligence (optional)	validate NPU, latency, skeleton
Full (finale)	all	none	the demo

This is the structural guarantee that **no single track's slippage sinks the demo**.

8. Part 1 — Pose abstraction layer & feature extraction

8.1 The seam (the single most important interface)

Per frame, the native side returns **one flat primitive array** to Kotlin:

```
float[17*3 + 5] = [ x0,y0,c0,  x1,y1,c1, ..., x16,y16,c16,  backendCode, detMs, poseMs, srcW, srcH ]
                  |_____| 51 keypoint floats (source-frame px) |_____| 5 metadata floats
```

- **No per-frame JNI object allocation** (no `NewObjectArray`). One primitive array, reused.
- `backendCode`: 0=CPU,1=GPU,2=NPU,3=MOCK — the backend *actually used*, read honestly from SNPE.
- Above the seam: nothing knows about tensors/DLCs. Below it: nothing knows about reps/exercises.

- **Consequence:** the entire intelligence layer is JVM-unit-testable with recorded keypoint CSVs as fixtures — no device, no NPU, no camera.

8.2 The 17 COCO keypoint indices (global constants, fix once)

```
0 nose · 1 L-eye · 2 R-eye · 3 L-ear · 4 R-ear ·
5 L-shoulder · 6 R-shoulder · 7 L-elbow · 8 R-elbow · 9 L-wrist · 10 R-wrist ·
11 L-hip · 12 R-hip · 13 L-knee · 14 R-knee · 15 L-ankle · 16 R-ankle
```

Skeleton bones for rendering: (5,7)(7,9)(6,8)(8,10)(5,6)(11,12)(5,11)(6,12)(11,13)(13,15)(12,14)(14,16)(0,5)(0,6).

8.3 Normalization (camera/scale/translation invariant) — `RealFeatureExtractor`

A **pure function** `normalize(PoseFrame) → NormalizedPose`. Steps:

1. **Confidence gate first.** For each keypoint, if `conf < τ_kp` (start **0.3**, tune on-device), mark it unusable. Any feature that needs an ungated keypoint emits `NaN`; downstream rules **skip**, never fabricate. (Fabricating values from low-conf joints is the #1 cause of phantom reps and wrong cues.)
2. **Hip-center** = midpoint(L-hip 11, R-hip 12). Translate all points so hip-center = origin.
3. **Torso length** = distance(hip-center, shoulder-center = mid(5,6)). Scale all coords by $1/\text{torso-length}$ → body-size and camera-distance invariant.
4. **Joint angles** (the primary features — translation/scale invariant and largely viewpoint-robust). Angle at B for triple (A,B,C) = angle between vectors BA and BC. Computed every frame:
 - `kneeL = ∠(hipL, kneeL, ankleL)`, `kneeR = ∠(hipR, kneeR, ankleR)`
 - `hipL = ∠(shoulderL, hipL, kneeL)`, `hipR = ∠(shoulderR, hipR, kneeR)`
 - `elbowL = ∠(shoulderL, elbowL, wristL)`, `elbowR = ∠(shoulderR, elbowR, wristR)`
 - `shoulderL = ∠(hipL, shoulderL, elbowL)`, `shoulderR = ∠(hipR, shoulderR, elbowR)`
 - `torsoLeanFromVertical` = angle of (shoulder-center → hip-center) vs image vertical
 - `openness` (jumping jack composite, range ~0=closed ... ~1=open): `armRaise = clamp01((shoulderCenter.y - wristCenter.y) / torsoLen)` (positive when wrists above shoulders; image y grows downward so a raised wrist has smaller y); `legSpread = clamp01((ankleSpread / shoulderWidth - 1.0) / 1.0)` (0 at feet-together ≈ shoulder width, 1 at ~2× spread); `openness = 0.5·armRaise + 0.5·legSpread`. All distances in torso-normalized units.
5. **Smoothing.** Each angle passes through an **EMA ($\alpha \approx 0.3$)** or 5-tap moving average before any threshold compare. Stateful, per-angle. This is what kills jitter-induced double counts.

Output `NormalizedPose{ jointAngles[], keypointsNorm[], confidences[], visible }`.

9. Part 1 — Exercise recognition & "is-exercising" gate

Two-tier, behind `ExerciseClassifier`, consuming a short ring-buffer window of `NormalizedPose`.

9.1 "Is-exercising" gate (satisfies P1.1)

Per frame, compute **motion energy** `E` = mean, over the candidate primary angles {knee, elbow, openness}, of each angle's **variance across the last 30 frames** (≈ 1 s @30 fps). State machine with hysteresis + dwell so it doesn't flicker between idle and active:


```
state IDLE:      enter EXERCISING when E > E_active for ≥ 0.5 s  (start E_active ≈ (15°)²)
state EXERCISING: return to IDLE      when E < E_idle   for ≥ 0.5 s  (start E_idle   ≈ (8°)²)
```

While **IDLE** the classifier emits **NONE** and the **RepCounter** is frozen — so we never count reps while a person stands still, fidgets, or walks into frame. Thresholds are tunable constants (= test vectors), tuned on recorded idle-vs-active clips via **tools/threshold_tuner**. This is a distinct, demoable feature: the HUD shows "Ready"/"Resting" vs the live exercise label.

9.2 RuleExerciseClassifier (guaranteed baseline, zero training)

Classify the Core 5 from the **signature** of the windowed features: (a) which primary angle has the largest oscillation amplitude, (b) torso orientation, (c) left-right symmetry. Decisive signatures:

Exercise	Signature	Primary angle	Discriminator
PUSHUP	torso ~horizontal, elbows oscillating, hips level with shoulders	elbow (avg)	horizontal torso vs curl
SQUAT	torso upright, both knees oscillating in sync, hip-center descends	knee (avg)	symmetric knees vs lunge
LUNGE	torso upright, asymmetric knee bend, one ankle forward	front-knee	asymmetry vs squat
BICEP_CURL	torso upright, shoulder fixed, elbow oscillating, no leg motion	elbow (active)	fixed shoulder + vertical vs pushup
JUMPING_JACK	wrists rise above head AND ankles spread, oscillating together	openness composite	both-limb openness

Hysteresis on the classification itself: don't switch the displayed exercise unless the new type wins for ~10 consecutive frames. Emit **UNKNOWN** when no signature is confident — **honesty over a wrong guess**.

9.3 LearnedExerciseClassifier (optional, additive — the scale path)

Pinned architecture (so it's buildable, not hand-wavy): a **1D temporal CNN (TCN-style)** over a fixed **32-frame window** (~1 s @30 fps). **Per-frame input feature vector = 61 floats** = 10 joint angles + 34 normalized keypoint coords (17×2) + 17 confidences → window tensor **[1, 32, 61]**. Network: 3 dilated 1-D conv blocks (dilation 1/2/4, ~64 channels, kernel 3, causal) → global average pool → dense → **softmax over 6 classes** {NONE, SQUAT, PUSHUP, LUNGE, BICEP_CURL, JUMPING_JACK}. Size **<5 MB** (INT8/FP16 TFLite), runs on **CPU** every ~8 frames (NPU budget stays for pose). Trained on Fit3D / InfiniteRep / Kaggle-workout keypoints (extracted with the same 17-COCO convention) + self-recorded sessions. **Critically, it consumes the exact FeatureExtractor output** — the device feature vector *is* the model input → no train/serve skew (see §16, the parity contract). If the model file is absent the fusion (§9.4) is silently rules-only.

9.4 FusedExerciseClassifier (honest fusion + fallback)

Runs rules always; runs learned if loaded. If learned **confidence > τ_learned** (≈0.7) → use learned, **source=LEARNED**; else rules, **source=RULES**; if both agree → **source=FUSED** (boosted). If the learned model is absent/failed to load → silently rules-only. The displayed exercise always carries its honest **source**. **Rules are the baseline; the learned model is the upgrade — its absence is invisible.**

Known confusion pairs to watch in tuning: push-up vs plank-hold (motion energy distinguishes), squat vs lunge (symmetry), curl vs pushup (torso orientation).

10. Part 2 — On-device runtime, `.dlc` conversion & the CPU-vs-NPU demonstration

10.1 What the reference gives us (keep verbatim)

`VisionSolution4` runs **two models on the NPU via SNPE C++**: YoloNAS_SSD (320×320) person-detect → HRNet (256×192) → 17 keypoints. The C++ call flow (in `inference.cpp` / `inference_helper.cpp`):

```
loadContainerFromBuffer() → zdl::DLContainer::IDLContainer::open(buffer,size)
checkRuntime()           → zdl::SNPE::SNPEFactory::isRuntimeAvailable(rt) // falls back DSP→GPU→CPU
setBuilderOptions()      → SNPEBuilder.setPerformanceProfile(BURST)
                        .setExecutionPriorityHint(HIGH)
                        .setRuntimeProcessorOrder(runtimeList)
                        .setUseUserSuppliedBuffers(true)
                        .setPlatformConfig(useAdaptivePD:ON).build()
createInput/OutputBufferMap() → getInput/OutputTensorNames(); getInputOutputBufferAttributes(name)
createUserBuffer()        → getUserBufferFactory().createUserBuffer(...) // zero-copy, Float32
execute(inputMap, outputMap) // YoloNAS, then HRNet
```

Runtime is passed down as a char: `'C'` CPU / `'G'` GPU / `'D'` DSP(HTP). **The in-app runtime toggle already exists.** We build on it; we do not re-derive it.

10.2 `.dlc` conversion — two independent paths (both prepped)

- **Path A — generate (Docker):** `Generate_models/GenerateDLC.ipynb` inside the QAI RT Docker converts + quantizes the upstream models, emitting `Quant_yoloNas_s_320.dlc` + `hrnet_axis_int8.dlc`. Needs a COCO calibration path (val2017 ~1 GB suffices). **Per-device:** re-run targeting the venue Hexagon arch.
- **Path B — download (AI Hub):** `aihub.qualcomm.com` → Computer Vision → Pose Estimation → pick the device (e.g. SM8750 or QCS6490) → download an already-INT8-quantized `.dlc` validated on that chip. **No Docker / Linux / COCO required.** This is our escape hatch (see §22 L3/L4/L6).
 - **Which model:** `HRNetPose-Quantized` (17 kpt) is the **drop-in** — same keypoint convention as the reference, so our `FeatureExtractor` and the seam are unchanged; pair with a person-detector or run it full-frame for the single-person coaching case. `MediaPipe-Pose` (33 kpt, single-stage) is the richer *upgrade* (more landmarks for form, drops the separate detector → easier `<80 ms`), but requires a **33→17 mapping** in `PoseEngine` (MP shoulders/elbows/wrists/hips/knees/ankles/nose/eyes/ears → COCO indices) and adapting pre/post-processing (PDF Ch 6.4). Default to HRNetPose-Quantized; treat MediaPipe-Pose as the latency/quality lever (§11) behind the same `PoseEngine` interface so nothing downstream changes.

10.3 The honest CPU-vs-NPU measurement (`Benchmark`)

1. **Warm-up discarded.** First ~10–30 inferences (or ~1 s) after a backend switch are excluded (SNPE graph prep / DSP spin-up are not steady state).
2. **Time the right thing.** Per-inference compute latency captured in C++ around the `execute()` calls (`detMs` , `poseMs`) — excludes JNI/Java jitter. Also end-to-end frame latency and pipeline FPS. All three displayed, **labelled distinctly** (judges conflate "inference ms" and "FPS"; we don't).
3. **Aggregate honestly.** Report **median (p50) and p90** over the last ~100 post-warm-up frames + an EMA for the live readout. Median, not mean (robust to GC/scheduler spikes).

4. **Same input, model, resolution** for both backends — only the runtime varies.
5. **Report the backend actually reached, never the requested one.** If `forceBackend(NPU)` falls back, the panel says "NPU requested → ran on CPU (DSP unavailable)."

10.4 The live panel (the Part-2 money-shot)

```
ENGINE
Backend: NPU (Hexagon HTP, V__)          ← honest; shows fallback if it happened
Pose infer: __ ms (p50) / __ (p90)        ← det + pose, measured in C++
End-to-end: __ ms      FPS: __
[ AUTO ] [ NPU ] [ GPU ] [ CPU ]          ← live toggle (reuse reference switch)
— A/B — NPU __ ms vs CPU __ ms
      speedup: __x                        ← computed live after a sweep
```

On stage: flip NPU→CPU and the FPS/latency visibly change. **A result that changes in front of the judge is the proof the NPU is real.** Corroborate with an offline QAIRT/AI-Hub profile screenshot on the slide.

10.5 The HTP-access reality (stated plainly, not buried)

The reference's HTP path needs `adb disable-verity; adb reboot; adb root; adb remount; adb shell setenforce 0` (SELinux permissive; **resets every reboot**). This works on the **venue QIDK (userdebug/engineering build)** but **likely fails on a retail phone** (iQOO Z5 / Moto can't `adb root`). → **NPU proof is a venue-device thing.** On retail phones we validate the pipeline on CPU/GPU and the panel labels that honestly. See §22 L2 and §17.

11. Part 2 — Performance budget & profiling methodology (<80 ms/frame)

No invented numbers. Fill `<__>` from on-device profiling. The budget is a *decomposition* so we know where to look when we miss the target.

```
PER-FRAME BUDGET (target end-to-end ≤ 80 ms → ≥ ~12.5 fps; stretch 33 ms → 30 fps)
[A] Camera acquire + YUV→RGB ..... < __ ms (pre-alloc, zero GC)
[B] Preprocess YoloNAS (320, BGR→RGB, ÷255) . < __ ms (C++)
[C] YoloNAS inference (320×320) on HTP ..... < __ ms (C++ execute) ← measure, expected dominant
[D] NMS (0.20) ..... < __ ms (C++)
[E] Crop + affine-warp 256×192 + ImageNet norm < __ ms (C++ preprocess_pose)
[F] HRNet inference (256×192) on HTP ..... < __ ms (C++ execute) ← measure, expected dominant
[G] Heatmaps(17) → argmax → keypoints ..... < __ ms (C++)
— SEAM (JNI float[51+5]) —
[H] FeatureExtractor (normalize+angles+EMA) . < __ ms (Kotlin, tiny)
[I] Classifier (rules/frame; learned every K) < __ ms (Kotlin; learned ~few Hz)
[J] RepCounter + FormAnalyzer ..... < __ ms (Kotlin, tiny)
[K] Compose overlay redraw ..... < __ ms (UI thread, decoupled – off critical path)
```

Rules: the <80 ms target is for the inference pipeline [A]–[G]; [H]–[J] are sub-ms pure Kotlin on 17 points; [K] is decoupled. **Levers if we blow 80 ms (in order):** (1) confirm HTP is *actually* engaged (check `backend()`, not the label); (2) run pose every frame but person-detect every N frames (track the box between detections — a person doesn't teleport); (3) lower analysis resolution; (4) drop to a single-stage pose model (MediaPipe-Pose / AI Hub one-shot) to remove an entire inference. The two NPU inferences ([C],[F]) will dominate — that's where INT8 + HTP BURST + zero-copy buy the budget.

12. Part 3 — Rep counting & per-exercise analyzers

StateMachineRepCounter — one FSM per exercise, all sharing this shape, driven by the smoothed primary angle θ :

States: TOP $\rightarrow$ DESCENDING $\rightarrow$ BOTTOM $\rightarrow$ ASCENDING $\rightarrow$ (rep++) $\rightarrow$ TOP
Two thresholds with a GAP (hysteresis): θ_{top} (extended) and θ_{bottom} (flexed), $\theta_{\text{bottom}} < \theta_{\text{top}}$
A rep counts on the full BOTTOM $\rightarrow$ TOP completion — never on a single threshold crossing.

Per-exercise primary angle + **starting** thresholds (tunable constants = test vectors; tuned on recorded data, never shipped as gospel):

Exercise	Primary θ	θ_{top}	θ_{bottom}	Notes
Squat	avg(kneeL,kneeR)	$\sim 165^\circ$	$\sim 95^\circ$	require both knees; symmetry checked in form
Push-up	avg(elbowL,elbowR)	$\sim 160^\circ$	$\sim 95^\circ$	torso must stay horizontal during the rep
Lunge	front knee	$\sim 165^\circ$	$\sim 95^\circ$	"front" = the more-flexed knee at bottom
Bicep curl	elbow (active arm)	$\sim 155^\circ$	$\sim 50^\circ$	shoulder must stay fixed
Jumping jack	openness	low (closed)	high (open)	full cycle = closed $\rightarrow$ open $\rightarrow$ closed

Anti-jitter rules (what makes counting robust): hysteresis gap (never a single threshold); **min-rep-duration** (reject "reps" faster than ~ 0.3 s — physically impossible, it's noise); **confidence-freeze** (if the primary joint drops below τ_{kp} mid-rep, freeze the FSM, resume on return; optionally void the rep); **partial-rep flagging** (reversed before reaching θ_{bottom} $\rightarrow$ **partial=true**, counted in a separate tally). The current FSM **phase** is exposed so the UI can show "go lower" while still DESCENDING. **Scalability proof:** adding an exercise = adding one analyzer config — no retraining.

Worked squat trace (smoothed avg-knee θ , $\theta_{\text{top}}=165^\circ$, $\theta_{\text{bottom}}=95^\circ$):

θ : 170 $\rightarrow$ [TOP] (standing)
150,120,100 $\rightarrow$ [DESCENDING] (θ falling, not yet < 95)
88 $\rightarrow$ [BOTTOM] (θ crossed below 95 $\rightarrow$ real depth reached)
100,140 $\rightarrow$ [ASCENDING]
168 $\rightarrow$ [TOP] $\Rightarrow$ rep++ (θ crossed back above 165 $\rightarrow$ one full rep)
PARTIAL example: 170 $\rightarrow$ 150 $\rightarrow$ 110(reverses) $\rightarrow$ 160 $\Rightarrow$ never hit BOTTOM (<95) $\rightarrow$ partial=true, no clean count.
NOISE example: a 1-frame 170 $\rightarrow$ 90 $\rightarrow$ 170 spike is rejected by min-rep-duration (≥ 0.3 s) + EMA smoothing.

13. Part 3 — Form analysis, scoring & coaching cues

RuleFormAnalyzer runs **per rep** (over the TOP $\rightarrow$ BOTTOM $\rightarrow$ TOP window). Each exercise has a small rule set; each rule contributes to a 0–100 score and may raise a **Cue**.

- **Squat:** depth (knee $\theta \leq$ target) $\cdot$ knee valgus (knees caving: knee-x vs ankle-x/hip-x) $\cdot$ symmetry ($|kneeL - kneeR|$) $\cdot$ torso lean (excessive forward) $\cdot$ tempo.
- **Push-up:** depth (elbow $\theta \leq$ target) $\cdot$ **hip sag/pike** (hip-center offset from shoulder–ankle line — the signature push-up fault) $\cdot$ elbow symmetry $\cdot$ tempo.
- **Lunge:** front-knee depth $\sim 90^\circ$ $\cdot$ knee-not-past-toes (knee-x vs ankle-x) $\cdot$ torso upright $\cdot$ L/R balance.

- **Bicep curl:** full ROM (extension→curl) · **shoulder/elbow drift** ("stop swinging, isolate the bicep") · symmetry · tempo (no jerking).
- **Jumping jack:** full arm extension overhead · full leg spread · arm/leg sync · rhythm.

Scoring: per-rep score = weighted sum of normalized rule sub-scores (depth, ROM, symmetry, alignment, tempo), clamped 0–100; weights are named constants. **Cue policy:** show **the single lowest-scoring rule that crossed its "bad" threshold**, phrased as an actionable imperative ("go deeper", "knees out", "hips up"). At most 1–2 cues at once — coaching, not a wall of text. Optional TTS. **Session summary:** total reps, avg score, most common fault. **Honest limit:** single-camera 2D pose can't see everything (depth ambiguity, occlusion) — form is **guidance, not medical advice** (stated in the UI and §24).

14. Part 3 — Gamification & session model

Light and tasteful — a *hint*, not a game (so it adds impact without diluting the engineering story): **streak** (consecutive days/sessions), **personal best** (reps at $\geq X$ form), a per-session **form grade** (A/B/C from avg score), and a rep goal with a progress ring. Session model: `Session{ exercise, reps[], avgScore, topFault, durationMs, grade }` persisted locally (Room/DataStore) — **on device, no account**.

15. Part 3 — App platform, camera & UI

- **Host:** fork of `VisionSolution4` (native camera + C++/SNPE kept). Coach UI in **Jetpack Compose** hosted over the preview (transparent `ComposeView` with `setZOrderOnTop`, or `AndroidView` wrapping the preview). **Skeleton redrawn in a Compose Canvas** from the keypoints crossing the seam (collapses two render paths into one) — *verify it profiles clean at 30 fps before deleting the GL path*.
 - **Camera:** `CameraEngine` (CameraX `ImageAnalysis`, `YUV_420_888`, `STRATEGY_KEEP_ONLY_LATEST`, pre-allocated buffers, zero per-frame alloc — lifted from the Pramāṇa pattern). On a board (RB3/QCS6490) the camera is USB/MIPI and the display is HDMI — adapt here, late (frontend-last).
 - **Camera-ownership decision (resolved):** Kotlin `CameraEngine` **owns the camera and converts YUV→RGB into one pre-allocated direct ByteBuffer**, passed *down* into `PoseEngine.infer(CameraFrame)` (which hands the buffer pointer to native `nativeInfer`). This matches the testable Pramāṇa pattern and keeps the seam clean. **Caveat (verify against the fork in Phase 1):** the reference's `CameraFragment` may feed its native side directly; if rewiring that fights us, the fallback is to let the **native side keep ownership of the camera** and have `CameraEngine` be a thin lifecycle wrapper that carries no pixels — `PoseEngine` then exposes a `SharedFlow<PoseFrame>` instead of `infer(frame)`. Decide after reading the forked `CameraFragment.java`; either way the intelligence layer (which only sees `PoseFrame`) is unaffected.
 - **Overlay (projector-legible):** big rep counter, one cue line (large, color-coded), per-rep score, exercise label, the engine/latency panel. Designed to read across a room on a projector.
 - **Lifecycle/permissions:** CAMERA only; no INTERNET permission in the prod path.
-

16. Complete tech stack (every piece + why)

Layer	Choice	Why
Inference SDK	QAI RT v2.46.0.260424 / SNPE C++	what the reference uses; <code>.dlc</code> on CPU/GPU/HTP, runtime-switchable
Models	YoloNAS_SSD + HRNet (<code>.dlc</code> , INT8)	the reference's proven pose pipeline on the NPU
Model source	GenerateDLC.ipynb (Path A) + AI Hub (Path B)	two independent paths to a device-matched <code>.dlc</code>
Native build	NDK r26c + OpenCV 4.13.0 Android SDK	exact versions <code>resolveDependencies.sh</code> expects
App language	Kotlin	our layer; reference Java kept where it works
UI	Jetpack Compose	fast, polished coach UI; Canvas overlay
Camera	CameraX	YUV zero-alloc analyzer; lifecycle-aware
Learned head (optional)	PyTorch → TFLite (CPU)	tiny temporal classifier; CPU keeps NPU budget for pose
Concurrency	Kotlin coroutines + StateFlow	one analyzer thread → immutable CoachState to UI
Persistence	DataStore/Room	local session history; no cloud
Build	Gradle + version catalog , AS Panda 4 v2025.3.4	reproducible; the AS version the PDF mandates
Profiling	adb / Android Studio profiler / QAI RT profile	the only valid source of latency numbers

The parity contract: `mL/features/` (Python) and the Kotlin `FeatureExtractor` implement the *same* normalization+angle math, pinned by parity tests — so the learned model has no train/serve skew.

17. Redundancy & failover (every failure has a pre-decided action)

Failure	Pre-decided action
NPU unreachable (no root / wrong arch)	Auto-fall to GPU/CPU; panel labels it honestly; demo runs on CPU; show prep NPU clip
Root denied on device	Full coaching demo on CPU + pre-recorded NPU-mode clip side-by-side
<code>.dlc</code> won't load / arch mismatch	Swap <code>PoseEngine</code> → <code>MockPoseEngine</code> (recorded keypoints); full coaching still demos
Camera/lighting bad	"Step into the box" floor marker; else pre-recorded full-demo clip, narrated
Total app freeze	3-min backup video (NPU mode), calm narration from the point of failure
Learned model bad/absent	Silently rules-only (baseline); nothing visible breaks

Meta-rule (from Pramāṇa): mock/canned values are **stage-rescue insurance, never the development default**. If you reach for them while building, the real thing is broken — fix it.

18. Module ownership & 3 parallel tracks

Three independent tracks; mock-everything + frozen interfaces keep them unblocked; collapses cleanly to fewer people.

- **Engineer A — "Metal" (device · pose · NPU):** Part 1 on-device + all of Part 2. Owns: fork bring-up; keep-verbatim C++/JNI/SNPE core; the seam (`float[51+5]`); `SnpePoseEngine` ; DLC gen (Path A) / AI Hub (Path B); device prep (root/ `setenforce` /skels); `Benchmark` + honest CPU-vs-NPU. Branches: `track/metal` .
- **Engineer B — "Brain" (exercise intelligence):** Part 1 recognition + Part 3 logic. Owns: `FeatureExtractor` , `RuleExerciseClassifier` , `StateMachineRepCounter` , `RuleFormAnalyzer` , `FusedExerciseClassifier` , the optional `ml/` learned head + dataset/training + `tools/threshold_tuner` . Works **100% device-free** against `MockPoseEngine` + recorded CSVs. Branches: `track/brain` .
- **Engineer C — "Glue → App" (integration, then frontend LAST):** integration + Part 3 app. Owns: `CoachOrchestrator` , `CameraEngine` lifecycle, mock-fallback DI (`App.kt`), CI, repo scaffold, `tools/keypoint_recorder` , and **the Compose UI in the final phase**. Branches: `track/app` .

Integration through interfaces only. Freeze `api/` in Phase 0; nobody edits another track's package.

19. Vibe-code vs Manual matrix

The hard "manual" surface here is **SNPE C++ + JNI + DLC + CMake/NDK** (not crypto, as in Pramāṇa).

MANUAL — do NOT vibe-code; keep reference verbatim; review after, don't "improve": SNPE/QAIRT C++ init & call flow · JNI marshalling / the seam (no per-frame objects) · DLC conversion+quantization (`GenerateDLC.ipynb` , QAIRT Docker, per-device regen) · CMake/NDK r26c + OpenCV linking (`resolveDependencies.sh`) · skel-lib bundling per Hexagon arch + `setenforce` device prep · affine-warp + ImageNet-norm preprocessing + NMS (keep verbatim — silent-failure-prone, already correct).

VIBE-with-verification — draft, then check against current docs / on-device: CameraX YUV zero-alloc · Compose + `AndroidView` / `SurfaceView` interop + z-order · Compose Canvas skeleton · Gradle/NDK/version-catalog/abiFilters · the exercise *threshold tuning* (verify against recorded data) · MediaPipe-Pose / AI-Hub model swap (verify the 33-vs-17 keypoint mapping).

VIBE-safe — drive freely, spot-check: the entire pure-Kotlin intelligence layer (`FeatureExtractor`, rule classifier, rep FSM, form analyzer) · all mocks · data classes/enums · unit tests · the PyTorch classifier + dataset prep + export · `Benchmark` aggregation math · docs.

Anti-patterns: don't let an assistant rewrite the SNPE/JNI/CMake core "to clean it up" (it's debugged — touching it loses a day); don't invent latency numbers; don't allocate / marshal objects per frame; don't fabricate values for low-confidence joints; don't claim NPU on a fallback.

20. Engineering rules (the non-negotiables)

1. **Zero per-frame allocation** — pre-allocate all buffers (RGB, the 51-float receiver, angle arrays, EMA ring buffers) at startup; reuse. No per-frame JNI object marshalling.
2. **Threading** — one dedicated single-thread analyzer executor; `STRATEGY_KEEP_ONLY_LATEST` ; the whole `infer→features→classify→reps→form` chain runs there and publishes an **immutable** `CoachState` to a

`StateFlow`. The **UI thread only renders** the latest `CoachState`; it never blocks the analyzer. Throttle person-detect + the learned classifier to every Nth frame; cache between.

3. **Confidence gating everywhere** — one `τ_kp`; skip, never fabricate; freeze the FSM on confidence loss.
4. **Honest reporting always** — `backend()` = the tier reached; panel labels requested-vs-reached; exercise carries its `source`; `UNKNOWN` over a wrong guess; partial reps flagged.
5. **The seam is sacred** — keypoints cross C++ → Kotlin as one flat primitive array at one boundary.
6. **Mock-fallback DI** — probe for the real dependency (DLC asset present? SNPE init ok?), fall back to the mock with a logged warning, per interface. The app always boots to *something* demoable.
7. **Tuned constants are test vectors** — every threshold lives in named config, validated against recorded sessions in CI. Drift = test failure.
8. **No cloud / no network in the hot path.** Privacy + the PS's "fully on-device" requirement.

21. Demo script & backup plans

~4 min live + buffer. Rehearse twice on the actual device. The backend badge always shows the truth.

- **0:00–0:30 — Honest open (inoculation).** "Qualcomm gave every team the same kit — pose on the NPU with a runtime toggle. So we won't pretend pose-on-NPU is our innovation. We built the layer it stops at." (One slide: the given-vs-built ledger. Buys immunity to the killer question.)
- **0:30–1:30 — Live coaching.** Do squats. Skeleton overlay, auto-detected "SQUAT", rep counter ticking, per-rep score flashing. "Detected, counting, scoring — all on the device, no cloud."
- **1:30–2:15 — Emotional wow: a bad rep gets corrected.** Do a shallow squat / knees caving. Score drops, one cue appears (spoken): "Go deeper — hips below knees." "One fix at a time, like a real coach."
- **2:15–3:15 — Technical wow: live CPU↔NPU.** Toggle NPU→CPU. Latency/FPS visibly change. "Same model, same app — Hexagon NPU vs CPU, measured live, real numbers from this device." Flip back; smoothness returns.
- **3:15–3:45 — Summary + scale tease.** Session summary (reps, avg form, top fault, streak). "Five exercises today; adding a sixth is one module; the hard part — pose on the NPU — Qualcomm solved, we built the coach on top."
- **3:45–4:00 — Close.** "Private, offline, on-device form coaching — the half of this problem the reference leaves to you. That's Vyāyāma."

Backups (each with a pre-decided trigger + rehearsed transition): (1) no root/HTP → full demo on CPU + pre-recorded NPU clip side-by-side; (2) model won't load → `MockPoseEngine` drives the full coaching stack on recorded keypoints; (3) camera/lighting fails → floor marker, else narrated backup video; (4) total freeze → 3-min backup video, calm narration.

Demo-day kit checklist (pack/verify before walking up):

- [] Venue device: charged, `adb root` + `setenforce 0` done (re-run after any reboot), APK installed, correct-arch DLC+skels in place, backend badge reads NPU
- [] Laptop: backup video (NPU mode, recorded on this device class) + slides (lead with the given-vs-built ledger) + the offline QAIRT/AI-Hub profile screenshot
- [] Second phone: side-by-side NPU-clip for the no-root fallback
- [] USB cable + a tripod/stand + floor marker tape (deterministic framing)

- [] Device in airplane mode (proves "no cloud" live) + `MockPoseEngine` build ready as the model-won't-load fallback
- [] Two rehearsals done, including each backup transition

22. Risk register & hidden-catch landmines

The real ones. Each: **what** · **when it bites** · **find it EARLY** · **defuse**.

#	Landmine	When it bites	Find early	Defuse
L1	RESOLVED — device confirmed: QIDK / Snapdragon 8 Elite (V79) / Android (organizer erratum on PS)	—	confirmed	Target the 8-Elite/V79 build: the reference <code>VisionSolution4</code> is validated on 8 Elite → its <code>.dlc</code> + skel <code>.so</code> are the right ones. Device-agnostic ladder retained as insurance.
L2	HTP needs <code>adb root</code> + <code>setenforce 0</code> (userdebug only)	when proving NPU	confirm at workshop	QIDK is a developer kit → almost certainly userdebug/rootable (unlike a locked retail phone); honest DSP→GPU→CPU fallback if not
L3	DLC gen needs Ubuntu 22.04 + Docker (18 GB); you're on Windows	model generation	Day-1 env setup	WSL2 Ubuntu; OR AI Hub Path B (no Docker/Linux/COCO)
L4	Wrong-arch DLC/ <code>.so</code> silently falls back to CPU	NPU validation	first on-device run	target V79 (8 Elite): use the reference's 8-Elite <code>.dlc</code> + skels, or an AI Hub 8-Elite DLC; <code>backend()</code> reports truth
L5	<80 ms/frame latency budget	latency claim	first on-device profile	trivial on V79 (8 Elite = top-tier NPU; HRNet pose ≈ low-single-digit ms); detect-every-N + single-stage pose are spare headroom; measured-only
L6	INT8 quant needs COCO calibration (GBs)	DLC gen	Day-1 (Docker route)	COCO val2017 subset (~1 GB); or Path B (already quantized)
L7	Fork build finicky (AS Panda 4 2025.3.4, OpenCV 4.13.0, NDK r26c, aar path)	bring-up	Phase 1	get UNMODIFIED reference running before any of our code ; if it fights >1 day, pivot to Path B + thin SNPE runner behind same <code>PoseEngine</code>
L8	Device-free dev needs recorded keypoint fixtures that don't exist	Phase 1 (B blocked)	it's a Phase-0 task	record Core-5 sessions Day 1 (reference dump / MediaPipe on laptop) → CSV
L9	Train/serve skew (Python features ≠ Kotlin features)	learned-model accuracy	parity tests in CI	one feature def mirrored both sides + parity tests; learned head is optional/additive
L10	3-person merge hell	integration	Phase 0	freeze <code>api/</code> ; per-track packages/branches; integrate via interfaces only; mock-everything

Master de-risks (structural): (a) **AI Hub Path B** neutralizes L3/L4/L6, de-risks L7. (b) **Rules-first** neutralizes L9. (c) **Mock-everything + frozen interfaces** neutralizes L10 and keeps the app demoable on partial completion. (d) **Honest `backend()` + measured-only numbers** neutralizes the "NPU claim collapses under questioning" risk. (e) **Prove the NPU path on ≥2 different Hexagon versions in prep** (a rooted personal Snapdragon + the

workshop QIDK 8-Elite/V79) → high confidence, plus an NPU-mode backup video recorded on the closest device class (8 Gen 2/3 ≈ V73/V75).

23. Timeline — frontend-LAST phased plan

Discipline: anything that *can* be done without the device *must* be done before it; only device-bound work waits for the kit.

Pre-workshop prep (Day 1 — env + reference + scaffold): WSL2 Ubuntu 22.04 + Docker (or commit to Path B); clone `quic/qidk`, `run_qairt_docker.sh`; get the **unmodified reference building** (and running on a personal phone, CPU/GPU); stand up the Vyāyāma repo — `api/` interfaces + domain types frozen, mock-fallback DI, **all mocks**, the **MockPoseEngine** + **recorded Core-5 CSV fixtures**, CI green.

Pre-workshop prep (Day 2 — the brain, device-free): `FeatureExtractor` (angles/normalize/smoothing) + unit tests; the rep FSM + per-exercise analyzers tuned on recorded data; form rules + scoring + one-cue engine + session summary + gamification skeleton; rule recognizer + is-exercising gate; (if ahead) start the learned head; `.dlc` conversion dry-run targeting the venue Hexagon class; start backup assets + deck.

Phasing across the build (frontend genuinely last): Phase 0 shared foundation → Phase 1 parallel core (A: reference on device · B: intelligence on mocks · C: orchestrator/DI/Benchmark on mocks) → Phase 2 real integration with a minimal debug overlay → **Phase 3 frontend** (the polished Compose UI) → Phase 4 polish + NPU backup video + rehearsal + deck with real numbers.

Device-gated (only at/after the workshop): confirm device → regenerate DLC + bundle correct-arch skels → root/ `setenforce` → real NPU validation + latency numbers → board camera/display adaptation if RB3/QCS6490. ~70–80% of the engineering is done device-free.

Hackathon-day hour budget (rough, ~18 h overnight build; adjust to the actual format):

Block	Hours	Work
Phase 1 — device bring-up	~3–4 h	confirm SoC/Hexagon; root + <code>setenforce</code> ; build/run reference on venue device; correct-arch DLC + skels
Phase 2 — real integration	~4–6 h	wire real <code>PoseEngine</code> through the seam; live keypoints; honest CPU-vs-NPU instrumentation; tune thresholds on real bodies
Phase 3 — frontend	~4–6 h	the polished Compose UI over real data; projector-legibility pass
Phase 4 — polish + demo	~3–4 h	NPU-mode backup video on venue device; deck with real numbers; rehearse twice incl. backups

Pre-hackathon effort estimate (the device-free build, done before the event): Phase-0 foundation ~½ day; intelligence layer + tests ~1–1.5 days; orchestrator/DI/Benchmark ~½ day; Compose UI (mock-driven) ~1 day; ml pipeline (optional) ~1 day; tools + CI + docs ~½ day. All parallelizable across the 3 tracks.

24. Honest limitations (what we explicitly do NOT claim)

- **2D single-camera pose has blind spots** — depth ambiguity, occlusion, extreme camera angles. Form feedback is **guidance, not medical advice**.

- **Rule-based form is per-exercise** — an honest scalability ceiling; the learned head is the scale path.
- **Recognition is the Core 5** — `UNKNOWN` for anything else, by design.
- **NPU proof is venue-device-bound** — retail phones (no root) run CPU/GPU; we label it honestly.
- **Latency numbers are device-specific** — only on-device profiling counts; the deck shows measured figures.
- **Single-person** — multi-person scenes are out of scope for v1.

On-screen honesty strings (the app says what it is): a one-time/footer line *"Form guidance, not medical advice."*; the backend badge shows the literal reached tier, and on fallback *"NPU requested → running on CPU (DSP unavailable on this device)"*; an `UNKNOWN` exercise shows *"Not sure what exercise — keep going"* rather than a guess; a low-confidence frame shows *"Step fully into frame"* rather than scoring garbage. These strings are not cosmetic — they are the honesty contract made visible to the judge.

25. Repo layout

```
docs/
  bible.md                this file - source of truth
  device-runbook.md       adb root / setenforce / DLC-regen / skel-bundling per device
  setup-checklist.md      QAIRT Docker / WSL2 / NDK r26c / OpenCV SDK / AS Panda 4 / fork bring-up
android/
  app/src/main/           Kotlin app (fork of VisionSolution4, repackaged)
    cpp/                  KEEP VERBATIM - inference.cpp · inference_helper.cpp ·
                          posedetectionYoloNAS.cpp · inc/zdl/ · CMakeLists.txt
    jniLibs/arm64-v8a/    libSNPE.so + skels V68/69/73/75(/79) [runbook; gitignored]
    assets/              *.dlc (per device) · recorded_sessions/*.csv (mock + test fixtures)
    kotlin/io/vyayama/
      api/                domain types (PoseFrame, NormalizedPose, ExerciseType, RepEvent,
                          FormFeedback, CoachState, Backend, PerfSnapshot) + the 7 interfaces
      pose/               SnpePoseEngine (JNI facade) · PoseJni.kt
      camera/             RealCameraEngine (CameraX YUV, zero-alloc)
      intelligence/feature/ RealFeatureExtractor
      intelligence/classify/ Rule/Learned/Fused ExerciseClassifier
      intelligence/rep/    StateMachineRepCounter (+ per-exercise configs)
      intelligence/form/   RuleFormAnalyzer (+ per-exercise rule cards)
      bench/              RealBenchmark
      coach/              CoachOrchestrator
      mocks/              MockPoseEngine (CSV replay) + one mock per interface
      ui/                 Compose: CameraScreen, overlay (Canvas), HUD, LatencyPanel,
                          ExercisePicker, gamification, Settings
      App.kt              hand-rolled DI + per-interface mock-fallback
      test/kotlin/        intelligence unit tests vs recorded_sessions/*.csv
ml/                       optional learned classifier + dataset tooling
  src/{data,features,models}, train.py, eval.py, export_tflite.py, reference_features.py
  tests/                  feature-parity tests (python vs pinned vectors)
tools/
  keypoint_recorder/      capture recorded_sessions CSVs from the live app
  threshold_tuner/        desktop harness: run rules over CSVs, print rep accuracy → tune constants
  .github/workflows/      android assemble + unit tests · python ml tests
CLAUDE.md · README.md
```

Appendix A — Interface contracts (copy-implementable)

```
// — domain types —
enum class Backend { AUTO, NPU, GPU, CPU, MOCK }
enum class ExerciseType { NONE, SQUAT, PUSHUP, LUNGE, BICEP_CURL, JUMPING_JACK, UNKNOWN }
enum class RepPhase { TOP, DESCENDING, BOTTOM, ASCENDING }

class PoseFrame(
    // crosses the seam
    val keypoints: FloatArray, // 51 = 17 * (x,y,conf), source-frame px
    val srcWidth: Int, val srcHeight: Int,
    val backend: Backend, // honest, what was reached
    val detLatencyMs: Float, val poseLatencyMs: Float,
    val timestampNs: Long
)
class NormalizedPose(
    val jointAngles: FloatArray, // fixed indexed set; NaN if ungated
    val keypointsNorm: FloatArray, // hip-centered, torso-normalized (x,y)
    val confidences: FloatArray,
    val visible: Boolean
)
class ExerciseState(val type: ExerciseType, val confidence: Float,
    val source: Source) { enum class Source { RULES, LEARNED, FUSED } }
class RepEvent(val exercise: ExerciseType, val index: Int, val phase: RepPhase,
    val completedNow: Boolean, val partial: Boolean, val formScore: Int)
class Cue(val severity: Int, val jointRef: String, val message: String)
class FormFeedback(val cues: List<Cue>, val perRepScore: Int,
    val symmetryPct: Float, val depthPct: Float, val tempoScore: Float)
class PerfSnapshot(val backend: Backend, val poseMsP50: Float, val poseMsP90: Float,
    val endToEndMs: Float, val fps: Float)
class CoachState(val pose: PoseFrame?, val normalized: NormalizedPose?,
    val exercise: ExerciseState, val rep: RepEvent?,
    val feedback: FormFeedback?, val perf: PerfSnapshot)
class CameraFrame( // produced by CameraEngine, consumed by PoseEngine
    val rgb: java.nio.ByteBuffer, // direct buffer, PRE-ALLOCATED + reused (never per-frame alloc)
    val width: Int, val height: Int, val timestampNs: Long)
class PoseWindow( // ring buffer of recent NormalizedPose, fed to the classifier
    val frames: ArrayDeque<NormalizedPose>, val capacity: Int) {
    fun push(p: NormalizedPose) { if (frames.size == capacity) frames.removeFirst(); frames.addLast(p) }
}
class AbResult(val npuMs: Float, val cpuMs: Float, val npuFps: Float, val cpuFps: Float,
    val speedup: Float, val backendReached: Backend)

// — interfaces (each has a real impl + a mock) —
interface PoseEngine {
    fun init(backend: Backend = Backend.AUTO): Backend // returns backend ACTUALLY reached
    suspend fun infer(frame: CameraFrame): PoseFrame
    fun backend(): Backend
    fun isReady(): Boolean
    fun forceBackend(b: Backend): Backend // re-init pinned; for live A/B
    fun lastPerf(): PerfSnapshot
    fun close()
}
interface CameraEngine { fun start(owner: LifecycleOwner); fun attachPreview(view: Any?)
    fun stop(); val frameBus: SharedFlow<CameraFrame>; fun close() }
interface FeatureExtractor { fun normalize(pose: PoseFrame): NormalizedPose }
interface ExerciseClassifier { fun classify(window: PoseWindow): ExerciseState; fun reset() }
interface RepCounter { fun update(ex: ExerciseType, p: NormalizedPose, tsNs: Long): RepEvent
    fun count(): Int; fun reset() }
interface FormAnalyzer { fun analyze(ex: ExerciseType, repWindow: List<NormalizedPose>,
    phase: RepPhase): FormFeedback; fun reset() }
interface Benchmark { fun onInference(b: Backend, detMs: Float, poseMs: Float, totalMs: Float)
    fun snapshot(): PerfSnapshot; fun startAB(); fun abResult(): AbResult?; fun reset() }

// — the native binding (the ONLY JNI surface; SnpePoseEngine wraps these) —
// Implemented in cpp/posedetectionYoloNAS.cpp (kept verbatim from the fork, signatures adapted).
external fun nativeInit(assetDir: String, runtime: Char): Int // 'C'/'G'/'D'; returns backend code reached
```

```
external fun nativeInfer(rgb: ByteBuffer, w: Int, h: Int, out: FloatArray): Int // fills out[51+5]; 0=ok
external fun nativeClose()
// 'out' is ONE pre-allocated FloatArray(56) reused every frame - no per-frame JNI object allocation.
```

Appendix B — Per-exercise rule cards (the biomechanics IP)

Each card: keypoints used · rep-FSM (primary θ , θ_{top} , θ_{bottom}) · form rules (priority order) · cues.

- **SQUAT** — kpts 5,6,11,12,13,14,15,16. FSM: θ =avg knee, top 165°, bottom 95°. Rules (priority): depth (knee $\leq$ 95°) → "go deeper, hips below knees"; valgus (knee-x inside ankle-x) → "knees out"; torso lean (>~45° from vertical) → "chest up"; symmetry ($|L-R| > 15^\circ$) → "even out your weight"; tempo.
- **PUSHUP** — kpts 5,6,7,8,9,10,11,12,15,16. FSM: θ =avg elbow, top 160°, bottom 95°; torso must stay horizontal. Rules: depth (elbow $\leq$ 95°) → "lower your chest"; hip sag/pike (hip off shoulder-ankle line) → "keep your hips in line"; elbow symmetry; tempo.
- **LUNGE** — kpts 11,12,13,14,15,16 + torso. FSM: θ =front knee, top 165°, bottom 95°. Rules: front-knee depth ~90°; knee past toes (knee-x beyond ankle-x) → "knee behind your toes"; torso upright; L/R balance.
- **BICEP_CURL** — kpts 5,6,7,8,9,10,11,12. FSM: θ =active elbow, top 155°, bottom 50°; shoulder fixed. Rules: ROM (full extend→curl) → "full range"; shoulder/elbow drift → "stop swinging"; symmetry; tempo.
- **JUMPING_JACK** — kpts 0,5,6,9,10,15,16. FSM: θ =openness composite, closed↔open. Rules: arm full overhead; leg full spread; arm/leg sync → "arms and legs together"; rhythm.

(Thresholds are starting values; tuned on recorded data via [tools/threshold_tuner](#) and pinned as test vectors. Filled with measured values in iteration 2–3.)

Appendix C — Glossary

SNPE Snapdragon Neural Processing Engine — high-level C++/Java inference API; loads [.dlc](#), dispatches to CPU/GPU/HTP. **QAIRT** Qualcomm AI Runtime — the unified SDK (SNPE + QNN). **QNN** AI Engine Direct — the low-level graph API. [.dlc](#) Deep Learning Container — Qualcomm's compiled model format. **HTP / Hexagon NPU** the fixed-function neural accelerator; runs INT8/INT16; lowest latency/power. **V68 / V79** Hexagon architecture versions (QCS6490 ≈ V68-class; SM8750/8-Elite = V79) — DLC/skel libs are arch-specific. **QIDK** Qualcomm Innovators Development Kit — a premium-Snapdragon dev board running Android. **RB3 Gen 2 / QCS6490** a robotics/IoT SBC (Hexagon 770, 12 TOPS, Android or Linux, HDMI + USB camera). **AIMET** AI Model Efficiency Toolkit (quantization). **AI Hub** aihub.qualcomm.com — catalogue of pre-converted [.dlc](#) models. **SELinux permissive / userdebug** the engineering-build state that allows app DSP access ([setenforce 0](#)). **COCO 17 keypoints** the standard body-joint set. **FSM** finite state machine. **EMA** exponential moving average. **Valgus** knees collapsing inward. **τ_{kp}** the keypoint-confidence gate threshold.

Appendix D — Self-audit log (the 3 brutal rating iterations)

Calibrated, lower-anchor-by-default scoring. Brutality rules applied: torn → lower; every score carries a gap; vague evidence capped at 75; an unsurvivable claim caps Honesty at 70; zero documented failure modes caps an aspect at 90. **Overall = mean of 9**. Target: >93 after iteration 3, no aspect <90.

#	Aspect	Iter 1 (v1)	Iter 2 (v2)	Iter 3 (v3)	Δ	What moved it across the iterations
1	Technical depth	80	90	92	+12	pinned learned head (TCN 32×61→6), JNI signatures, camera-ownership + fallback, openness + is-exercising formulas, worked rep trace
2	PS alignment	88	92	95	+7	is-exercising upgraded to a full FSM; §4 mapping complete; all 3 parts rigorous
3	Feasibility	84	85	93	+9	hackathon-day hour budget + pre-event effort estimates; device-free majority; mock-everything
4	Differentiation	90	90	94	+4	sharpened vs a concrete reskin baseline; the ledger + the Q&A weapon
5	Robustness/risk	88	90	94	+6	Path-B model choice, camera fallback, 2-device-proof de-risk, full failover table
6	Completeness	82	88	95	+13	CameraFrame/PoseWindow/AbResult/JNI defined; worked trace; this self-audit log; build checklist (App. E)
7	Demo readiness	87	87	94	+7	demo-day kit checklist added to the minute-by-minute + 4 backups
8	Clarity/actionability	86	89	95	+9	build-order checklist; pinned specs; copy-implementable interfaces; next-actions
9	Honesty	93	93	96	+3	on-screen honesty strings; measured-only enforced; every limit owned before a judge can raise it
Average		86.4	89.3	94.2	+7.8	iteration deltas +2.9 / +4.9 (calibrated — no jump >12); final >93, min aspect 92

Known residual gaps (honest, for a future v4): a 2nd worked rep trace (push-up hip-sag) and an explicit MediaPipe 33→17 keypoint mapping table would push Technical depth toward 96; a real prep day-by-day with named owners would push Feasibility toward 96. None are blocking — they are polish, not correctness.

Appendix E — Build-order checklist (execute top-to-bottom)

Phase 0 (shared) → 1 (parallel) → 2 (integration) → 3 (frontend) → 4 (polish). Tick as you go.

```
[ ] P0 Repo scaffold: Gradle + version catalog + manifest (CAMERA only, no INTERNET) + settings
[ ] P0 api/ - freeze all domain types + the 7 interfaces + the JNI externs (Appendix A)
[ ] P0 All mocks incl. MockPoseEngine (CSV replay) + mock-fallback DI (App.kt)
[ ] P0 Record Core-5 keypoint CSV fixtures (unblocks Brain track) ; CI green (assembleDebug + tests)
[ ] P1-A Fork quic/qidk VisionSolution4; build UNMODIFIED; run on a phone (CPU/GPU); confirm runtime toggle
[ ] P1-B FeatureExtractor (angles/normalize/EMA) + unit tests vs CSVs
[ ] P1-B RuleExerciseClassifier + is-exercising gate + RepCounter FSM + RuleFormAnalyzer + tests
[ ] P1-C CoachOrchestrator (StateFlow) + RealBenchmark (EMA/p50/p90/fps) on mocks
[ ] P2-A SnpePoseEngine: wire the seam (float[51+5]); real keypoints; honest backend(); forceBackend A/B
[ ] P2-B Tune thresholds on recorded data (threshold_tuner); finalize per-exercise rule cards
[ ] P2-B (optional) LearnedExerciseClassifier: train (ml/) → export TFLite → FusedExerciseClassifier
[ ] P2-C Integrate A+B behind the orchestrator with a minimal debug overlay; prove end-to-end data flow
[ ] P3-C Compose UI: skeleton Canvas overlay, HUD, rep counter, cues, latency A/B panel, gamification, summary
[ ] P4 Tune on real bodies; record NPU-mode backup video; deck with real numbers; rehearse twice + backups
```

Device-gated (only at/after the workshop): confirm SoC → correct-arch DLC + skels → root/ [setenforce](#) → NPU validation + real latency numbers → board (USB-cam/HDMI) adaptation if RB3/QCS6490.